

homework3

Yi-Ling Chen

May 17, 2021

```
## Loading required package: Matrix
## Loaded glmnet 4.1-1
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
## Loading required package: lattice
##
## Attaching package: 'MASS'
## The following object is masked from 'package:dplyr':
##
##   select
```

Task2.1

```
df <- read.table("prostate_data.txt", header=T)
sub_df <- df %>% filter(train==TRUE)
test_df <- df %>% filter(train==FALSE)

train_X = as.matrix(sub_df[1:8])
train_y = sub_df$lpsa
test_X = as.matrix(test_df[1:8])
test_y = test_df$lpsa

model1 <- lm(lpsa ~ lcavol, data=sub_df)
model2 <- lm(lpsa ~ lcp, data=sub_df)
model3 <- lm(lpsa ~ lcavol + lcp, data=sub_df)
MS1 <- summary(model1)$coef
MS2 <- summary(model2)$coef
MS3 <- summary(model3)$coef
print("Model1:")

## [1] "Model1:"
MS1
```

```
##           Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) 1.5163048 0.14772483 10.264387 3.123765e-15
## lcavol      0.7126351 0.08199036  8.691694 1.733134e-12
```

```
print("Model2:")
```

```
## [1] "Model2:"
```

```
MS2
```

```
##           Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) 2.5427013 0.13121183 19.378598 1.041939e-28
## lcp         0.4218252 0.09327981  4.522149 2.659180e-05
```

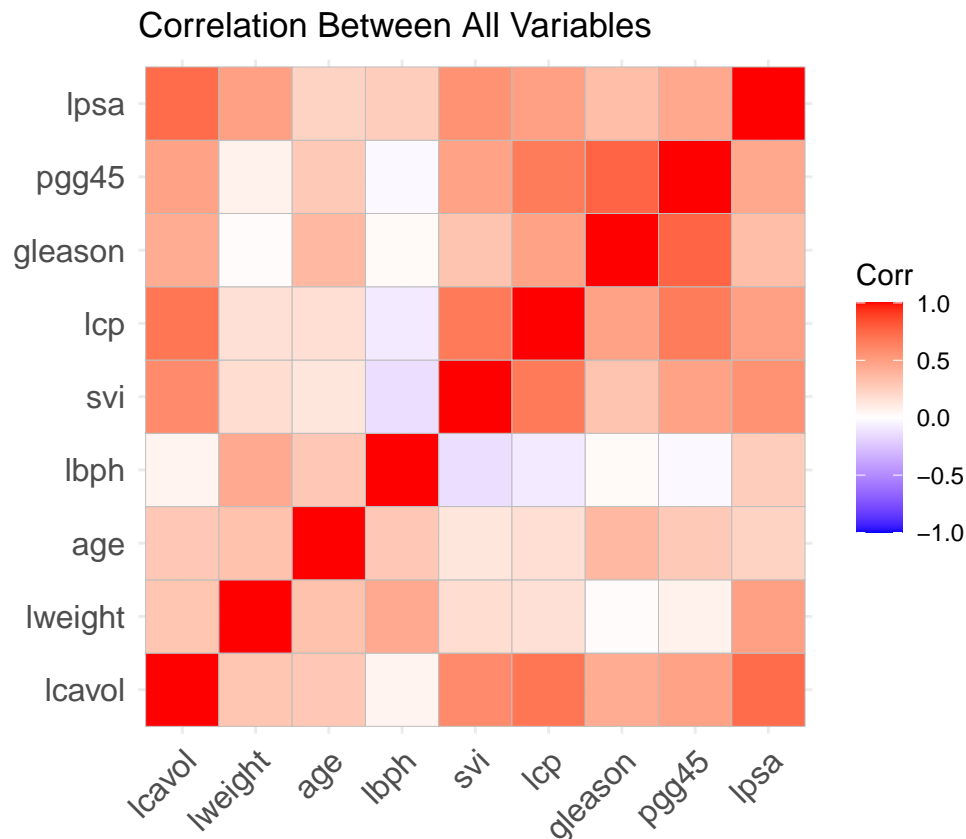
```
print("Model3:")
```

```
## [1] "Model3:"
```

```
MS3
```

```
##           Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) 1.47905119 0.1947748  7.5936468 1.678372e-10
## lcavol      0.73609361 0.1143882  6.4350510 1.802808e-08
## lcp        -0.03007034 0.1014736 -0.2963367 7.679325e-01
```

```
ggcorrplot(cor(sub_df[1:9]), title = "Correlation Between All Variables")
```



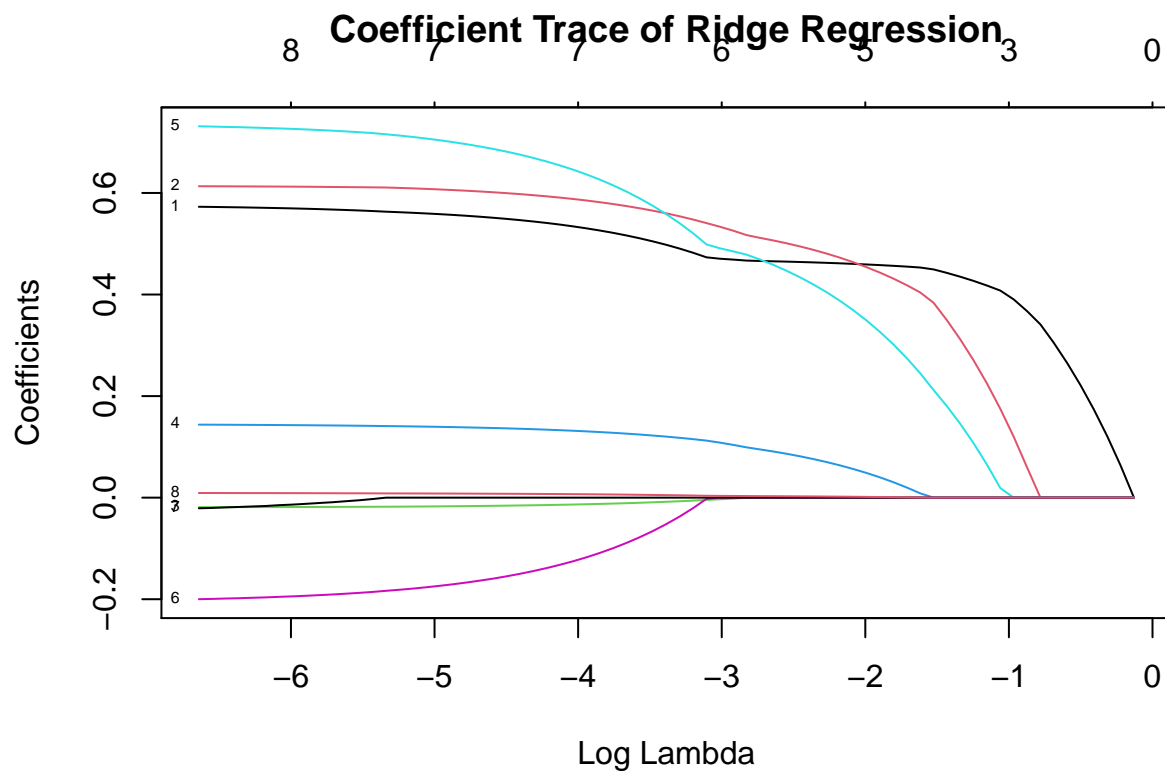
Since predictors lcavol(log cancer volume) and lcp(log of capsular penetration) are highly correlated. The more highly correlated the independent variables are, the more difficult to determine how much variation in Y each X is responsible for. Therefore, the standard errors for both variables become larger.

Task2.2

```
# least-square with all variables
least_square <- lm(lpsa~., data = df, subset = train)

# ridge regression
# tracing lambda estimation
ridge_trace <- glmnet(train_X, train_y, family = "gaussian", alphah=0)
# choosing best lambda with cross-validation
ridge_cv <- cv.glmnet(train_X, train_y, family = "gaussian", alphah=0)

plot(ridge_trace, label = TRUE, xvar = "lambda", main="Coefficient Trace of Ridge Regression")
```



```
coef(least_square)
```

```
## (Intercept)      lcavol      lweight      age      lbph      svi
## 0.429170133 0.576543185 0.614020004 -0.019001022 0.144848082 0.737208645
##      lcp      gleason      pgg45      trainTRUE
## -0.206324227 -0.029502884 0.009465162          NA
```

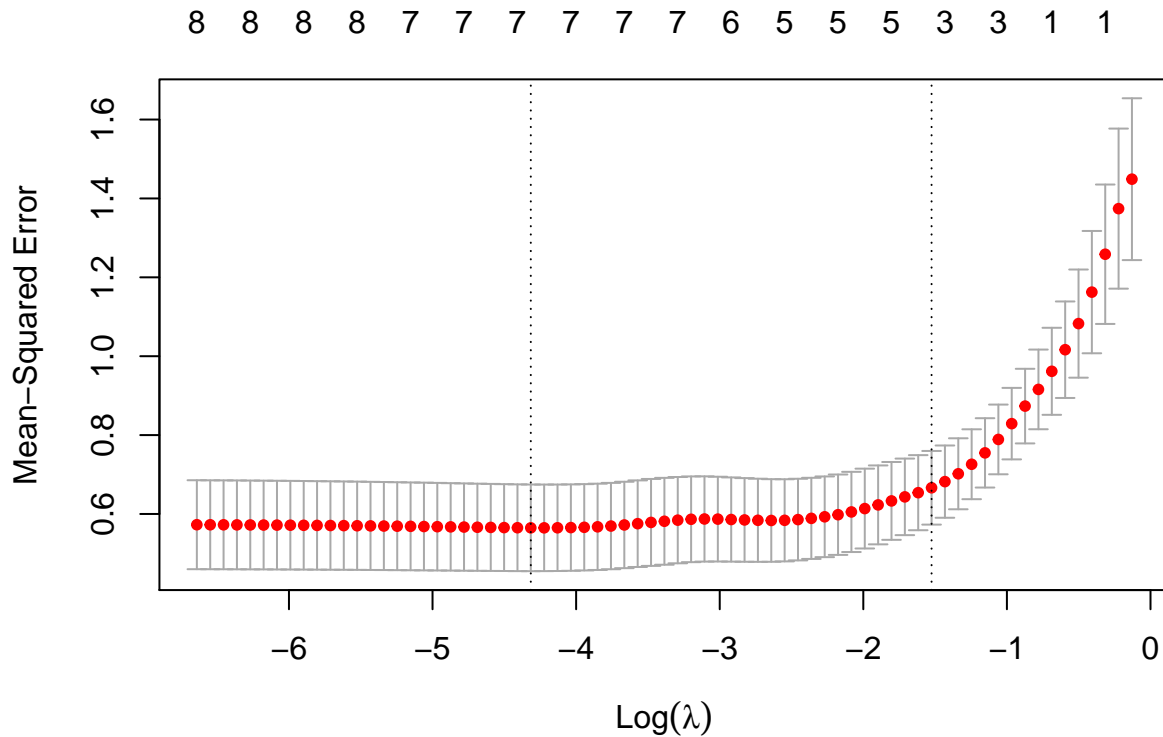
When lambda is zero, we can see the coefficients at the left side for each variables are the same as least-square model.

```
print(ridge_cv)
```

```
##
## Call:  cv.glmnet(x = train_X, y = train_y, family = "gaussian", alphah = 0)
##
```

```
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.01336   46  0.5649 0.10983       7
## 1se 0.21771   16  0.6664 0.09337       3
```

```
plot(ridge_cv)
```



According to `lambda.min` and `lambda.1se`, we can choose `lambda.min` for our model, which means when `lambda` is 0.01609, the cross-validation result has the lowest error. However, to the future unseen data, it's better to have a more regularized model because we don't know the distribution of the unseen data. Therefore, choosing the result of `lambda.1se` for the model.

```
# RMSE
best_ridge <- glmnet(train_X, train_y, family = "gaussian", alphah=0, lambda = ridge_cv$lambda.1se)
sqrt(mean((predict(best_ridge, test_X)-test_y)^2))
```

```
## [1] 0.6953873
```

```
min_ridge <- glmnet(train_X, train_y, family = "gaussian", alphah=0, lambda = ridge_cv$lambda.min)
sqrt(mean((predict(min_ridge, test_X)-test_y)^2))
```

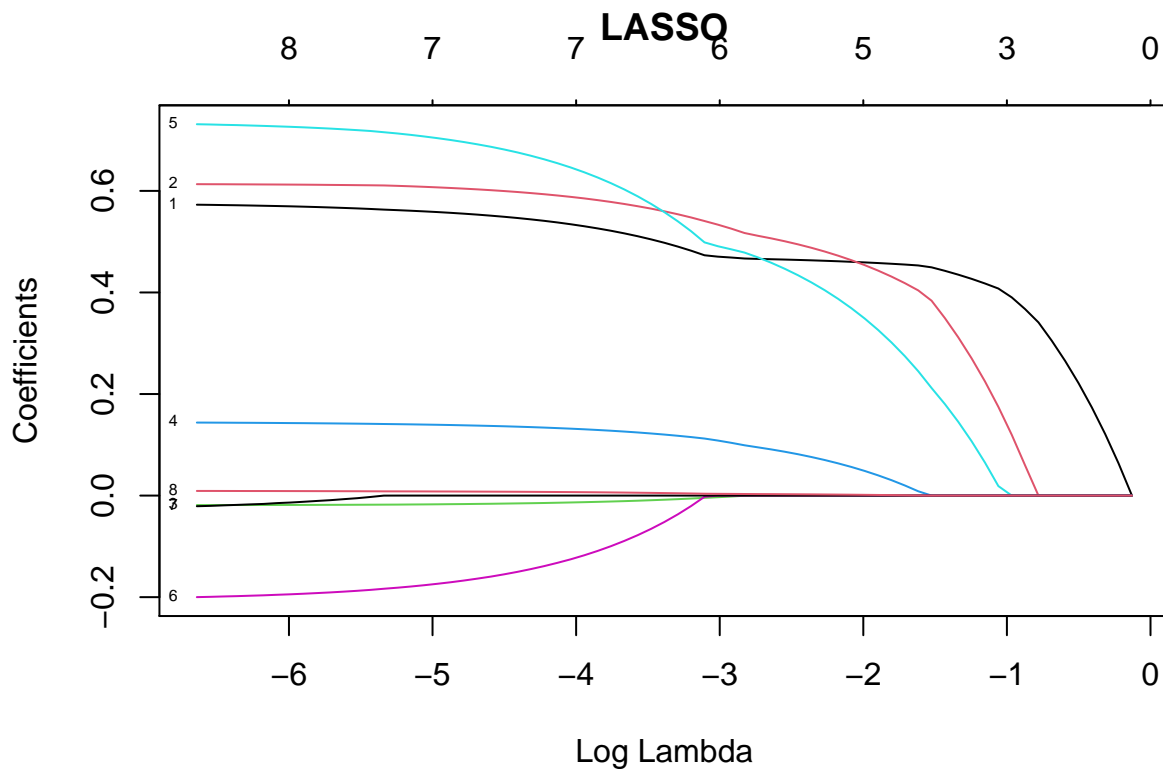
```
## [1] 0.7023585
```

The first value is the `lambda` from `lambda.1se`, which represents a more general model with less model complexity, the `rmse` is 0.6729088. The second one is the model with `lambda.min`, the `rmse` is 0.6993649. The result shows that the model `lambda` has a better performance.

Task2.3

```
# glmnet default alpha =1
lasso = cv.glmnet(train_X, train_y,family = "gaussian", alphah=1)
lasso_trace = glmnet(train_X, train_y,family = "gaussian", alphah=1)
print(lasso)

##
## Call: cv.glmnet(x = train_X, y = train_y, family = "gaussian", alphah = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.01466   45  0.5733 0.1199        7
## 1se 0.21771   16  0.6901 0.1016        3
plot(lasso_trace, label = TRUE,xvar = "lambda", main="LASSO")
```



```
lasso_1se = glmnet(train_X, train_y,family = "gaussian", alphah=1, lambda = lasso$lambda.1se)
print("The rmse estimate of lasso:")
```

Comparison of the rmse performance

```
## [1] "The rmse estimate of lasso:"
sqrt(mean((predict(lasso_1se, test_X)-test_y)^2))
```

```
## [1] 0.6953873
print("The rmse estimate of ridge regression:")

## [1] "The rmse estimate of ridge regression:"
sqrt(mean((predict(best_ridge, test_X)-test_y)^2))

## [1] 0.6953873
```

```
coef(lasso_1se)
```

LASSO coefficients

```
## 9 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept) 0.4228791
## lcavol      0.4494666
## lweight     0.3837892
## age         .
## lbph        .
## svi         0.2118731
## lcp         .
## gleason     .
## pgg45       .
```

```
coef(best_ridge)
```

Ridge coefficients

```
## 9 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept) 0.4228791
## lcavol      0.4494666
## lweight     0.3837892
## age         .
## lbph        .
## svi         0.2118731
## lcp         .
## gleason     .
## pgg45       .
```

After penalizing, both models has 5 predictors, all the predictors have the similar estimation. And LASSO has a lightly better result on rmse estimation.

Task2.4

Comparison of three models

```
# Search for the best elasticnet(combination of alpha and lambda)
for(i in 0:10){
  assign(paste("fit", i, sep=""), cv.glmnet(train_X, train_y, type.measure="mse", alpha=i/10,family="gaussian"))
}
```

```
# Show the best one
fit9
```

```
##
## Call: cv.glmnet(x = train_X, y = train_y, type.measure = "mse", alpha = i/10, family = "gaussian")
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.01352   47   0.618 0.1276         7
## 1se 0.26548   15   0.733 0.1711         3
```

```
alpha = 0.9
lambda = 0.18298
L1 = alpha*lambda
L2 = ((1-alpha)/2)*lambda
L1
```

```
## [1] 0.164682
```

```
L2
```

```
## [1] 0.009149
```

The elasticnet with $\alpha=0.3$, $\lambda_{1se}=0.31413$ has the best estimation of MSE on training data, and we can calculate the λ_{d1} and λ_{d2} according to the loss function. So $\lambda_{d1}(L1)=0.164682$, $\lambda_{d2}(L2)=0.009149$.

```
cv_ridge <- cv.glmnet(train_X, train_y, family = "gaussian", alphah=0)
cv_lasso <- cv.glmnet(train_X, train_y, family = "gaussian", alphah=1)
cv_elnet <- cv.glmnet(train_X, train_y, family = "gaussian", alphah=0.9)

cv_ridge
```

CV-RMSE: Ridge Regression

```
##
## Call: cv.glmnet(x = train_X, y = train_y, family = "gaussian", alphah = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.00527   56   0.5639 0.07711         7
## 1se 0.13673   21   0.6409 0.06600         5
```

```
cv_lasso
```

CV-RMSE: LASSO

```
##
## Call: cv.glmnet(x = train_X, y = train_y, family = "gaussian", alphah = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.00189   67   0.5531 0.09769         8
```

```
## 1se 0.18074    18  0.6505 0.08842      5
```

```
cv_elnnet
```

CV-RMSE: Elasticnet

```
##
## Call:  cv.glmnet(x = train_X, y = train_y, family = "gaussian", alphah = 0.9)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.0101    49  0.6151 0.1246        7
## 1se 0.2177    16  0.7281 0.1608        3
```

```
# ridge
sqrt(0.6462)
```

```
## [1] 0.8038657
```

```
# lasso
sqrt(0.6430)
```

```
## [1] 0.8018728
```

```
# elasticnet
sqrt(0.7039)
```

```
## [1] 0.8389875
```

According to the cv-rmse, LASSO is the best and then ridge regression, the worst one is elasticnet.